CONFIDENTIAL — PRE-READ

AI-Native Experimentation Platform

From idea to live customer experiment in under 60 seconds — using nothing but a chat window.

Please read before the demo session.

This document sets context so you can watch for the right moments during the live walkthrough.

THE PROBLEM WE SOLVED

Experiments Shouldn't Take Three Weeks

Today, running a customer experiment looks like this: a product manager writes a spec, an engineer deploys new code, a data scientist interprets the numbers, and three weeks later there's a meeting to decide if it worked.

> **We built something different.** A product manager types a plain-English instruction into a chat window. An AI creates the experiment, sets the targeting rule, and the live app immediately serves a different experience to the right customers. The product manager never writes code, never opens a deployment pipeline, and never needs an engineer in the room.

During this demo you will watch a product manager — played by the presenter — take a grocery app from *zero personalization* to *five live customer segments*, then run a full A/B statistical test and roll out the winner, entirely through conversation.

What You Will See

The demo runs on a simulated Safeway-style grocery app. Six simulated customers — each representing a real segment — visit the same URL. The server reads their attributes (loyalty tier, location, shopping frequency, household type) and assembles a completely different page for each one. No browser-side logic. No separate codebases. One app, infinite personalizations, all configured by chat.

DemoApp2 — anonymous visitor, default page (no flags set)

01a-before-grid.png

> **The 60-second proof:** Watch the timestamp from the first chat message to the moment the live app shows a different experience for a targeted user. In every demo scenario, it is under 60 seconds. No engineer was involved. No code was deployed.

WHO YOU'LL SEE

Six Customer Personas

Each persona represents a real customer segment. Their attributes drive every experiment. Switching personas in the app instantly shows what that segment experiences.

Name	Tier	Location	Shopping	Household
Sarah Chen	Gold	California	Frequent	Family
Marcus Johnson	Silver	New York	Occasional	Couple
Emily Rodriguez	Basic	Illinois	New	Single
Robert Williams	Gold	Oregon	Frequent	Senior
Jordan Lee	Silver	Texas	Frequent	Couple
Priya Patel	Gold	Washington	Frequent	Single

Why this matters: These six people visit the same URL. The server evaluates each flag against their attributes and assembles a unique page layout — in real time — for each one. There is no frontend code for each segment. There is no separate deployment. The product manager added a targeting rule and it was live instantly.

Sarah Chen (Gold) — premium hero banner

05c-gold-hero.png

Marcus Johnson (Silver) — standard carousel

05d-silver-hero.png

LIVE DURING THE DEMO

Five Personalization Experiments

Each experiment below is created live, through chat, during the demo. The presenter types one message; the AI calls the API; the app changes immediately.

1

Gold Member Hero Banner

"Gold members should feel recognized the moment they land."

`membership_tier = gold`

Gold members see a premium hero banner with a gold gradient overlay, loyalty points callout, and "Gold Member Exclusive" pill. Silver and Basic members see the standard carousel — they are never targeted.

2

Frequent Shopper Deals Grid

"Frequent shoppers scan fast. Give them density."

`shopping_behavior = frequent`

Frequent shoppers see weekly deals as a dense 2-column grid — more products visible at once. Occasional and new shoppers see a horizontal scroll strip — easier to browse without overwhelm.

3

Family Category Boost

"Families buy snacks. Surface snacks first."

`household_type = family`

Family households see "Snacks & Chips" as the first category tile — boosted because families with kids buy snacks frequently. Senior households see "Meat & Seafood" first. The server reorders 16 categories in real time.

4

Regional Price Sort

"California shoppers are price-sensitive."

`location = CA`

Weekly deals for California users are sorted lowest-price-first. New York and other states continue to see popularity-sorted deals. One flag, one rule — regional behavior without a single line of new code.

5 Senior Fresh Section

"Control what every senior sees — instantly, for everyone."

`household_type = senior`

The fresh produce section can be shown or hidden for senior users via a single chat message. Robert Williams (the only senior persona) toggles between seeing and not seeing the section — no page reload needed for the flag change to take effect.

Sarah Chen — dense deals grid (frequent shopper)

`06c-frequent-grid.png`

Marcus Johnson — horizontal scroll deals (occasional)

`06d-occasional-scroll.png`

DEMO 10 — STATISTICAL TESTING

A Full A/B Lifecycle in One Conversation

After the five targeting experiments, the demo shows that the same platform also handles rigorous A/B testing with statistical significance — all through the same chat interface.

The Conversation Flow

Turn	What the presenter types	What happens
1	"Create an A/B test for the deals layout — scroll vs. grid, 50/50 split."	AI calls the API. Experiment created. Traffic split active.
2	"Get the results for that experiment."	AI retrieves conversion rates, uplift %, p-value, and confidence interval. Grid wins at $p < 0.05$.
3	"Roll out the winning variant to everyone."	AI updates the targeting rule to 100% grid. All users immediately see the winner.
4	(Presenter signs in as three personas)	All three personas see the grid — the rollout is confirmed live on screen.

> **What this proves:** Create, measure, decide, and deploy — the full experiment lifecycle — in a single conversation. No data analyst needed to read the results. No engineer needed to push the rollout. The AI does all of it.

Control Room chat — A/B results with statistical significance

10c-ab-results.png

What You Should Watch For

? How long from first message to live change on screen?

In every experiment, it is under 60 seconds. Time it yourself.

? Who is involved?

Only the product manager. No engineer wrote any code during the demo. No deployment ran.

? What is the rollback story?

The presenter can send one more chat message to revert any flag. Watch the page change again.

? How confident can we be in A/B results?

The platform reports p-value, confidence interval, and absolute uplift — standard statistical outputs — directly in the chat response.

? Could this work on our actual product?

The architecture is stack-agnostic. Server-side evaluation + pure-renderer components + AI-mediated flag management applies to any modern web app.

HOW IT WORKS

The Architecture in Plain English

You don't need to understand the code to understand the architecture. Here is what happens each time the presenter types a message:

#	Step	Who/What does it
1	Product manager types a natural-language instruction	Control Room chat UI
2	AI parses the intent and calls the flag management API	Claude (Anthropic) + 13 custom tools
3	Flag is created or updated with the targeting rule	GrowthBook (our Experimentation Engine)
4	Customer signs in; attributes sent to Layout API	DemoApp2 browser → /api/demoapp2/layout
5	Layout API evaluates all flags server-side and returns a complete page layout as JSON	Node.js server (Next.js API route)
6	Browser renders exactly what the server returned — no experiment logic runs client-side	React (pure renderer)

> **Zero experiment logic in the browser.** Every decision about what a customer sees — which banner, which sort order, which categories, whether a section is visible — is made by the server. The browser is a pure renderer. This is why adding a new experiment requires no React code changes and no frontend deployment.

Demo Timing

Time	Segment
0:00	Open DemoApp2 — show anonymous default page
1:00	Demo 1: Global flag → instant layout change
3:00	Demo 5: Gold member hero banner
6:00	Demo 6: Frequent shopper deals grid
9:00	Demo 7: Family category boost
11:00	Demo 8: Regional price sort (California)
13:00	Demo 9: Senior fresh section

Time	Segment
16:00	Demo 10: A/B test → statistical results → rollout
19:00	Architecture overview + automated validation run
20:00	Q&A

The One Sentence to Remember

> "A product manager types a plain-English instruction. In under 60 seconds, every customer in that segment sees a different experience — with no code deployed, no engineer involved, and no page reload required."